

# BtrFS: A fault tolerant File system

Kumar Amit Mehta

Date: 12/06/201

Coursework:

IAF0530



# Agenda

- Motivation
- Background
- BtrFS features
- Main principles and technical details
- Test results
  - Robustness
  - Performance
- Demo
  - Fault tolerance
- Conclusion
- Q&A

# Motivation

- Relevance



Source: [dilbert.com](http://dilbert.com)

# Motivation contd...

- Main factors for Storage faults
  - Firmware/Software Bug
  - Hardware Failure
    - Aging
    - Faulty hardware
    - Annual disk replacement rates typically exceed 1%, with 2–4% common and up to 13% observed on some systems [7].
  - Complexity
    - SUN took 7 years to develop ZFS and another few years more until ZFS was deployed in mission critical systems.
    - “File System are hard” - Ted Ts'o (Maintainer of ext4)



# Motivation contd...

**“Talk is cheap. Show me the code.”**

Linus Torvalds

[#] PCbots Lab's




pcbots.blogspot.com

```
332      /* Sun, you just can't beat me, you just can't. Stop trying,  
333      * give up. I'm serious, I am going to kick the living shit  
334      * out of you, game over, lights out.  
335      */  
336      .align 8  
337      .globl __csum_partial_copy_sparc_generic
```

<Snip from linux/arch/sparc/lib/checksum\_32.S>

# Background

- Historical Perspective
    - “B-trees, Shadowing, and Clones Ohad Rodeh [1]” (USENIX, 2007)
    - Chris Mason (Combined ideas from ReiserFS and COW friendly B-trees as suggested by Rodeh )
    - Finally accepted in mainline Linux Kernel in 2009
    - Default root File system for SuSE, Oracle Linux
    - 2014, Facebook announced [2] to use BtrFS as Trial
- 

```
Date      Tue, 12 Jun 2007 12:10:29 -0400
From      Chris Mason <>
Subject   [ANNOUNCE] Btrfs: a copy on write, snapshotting FS
```

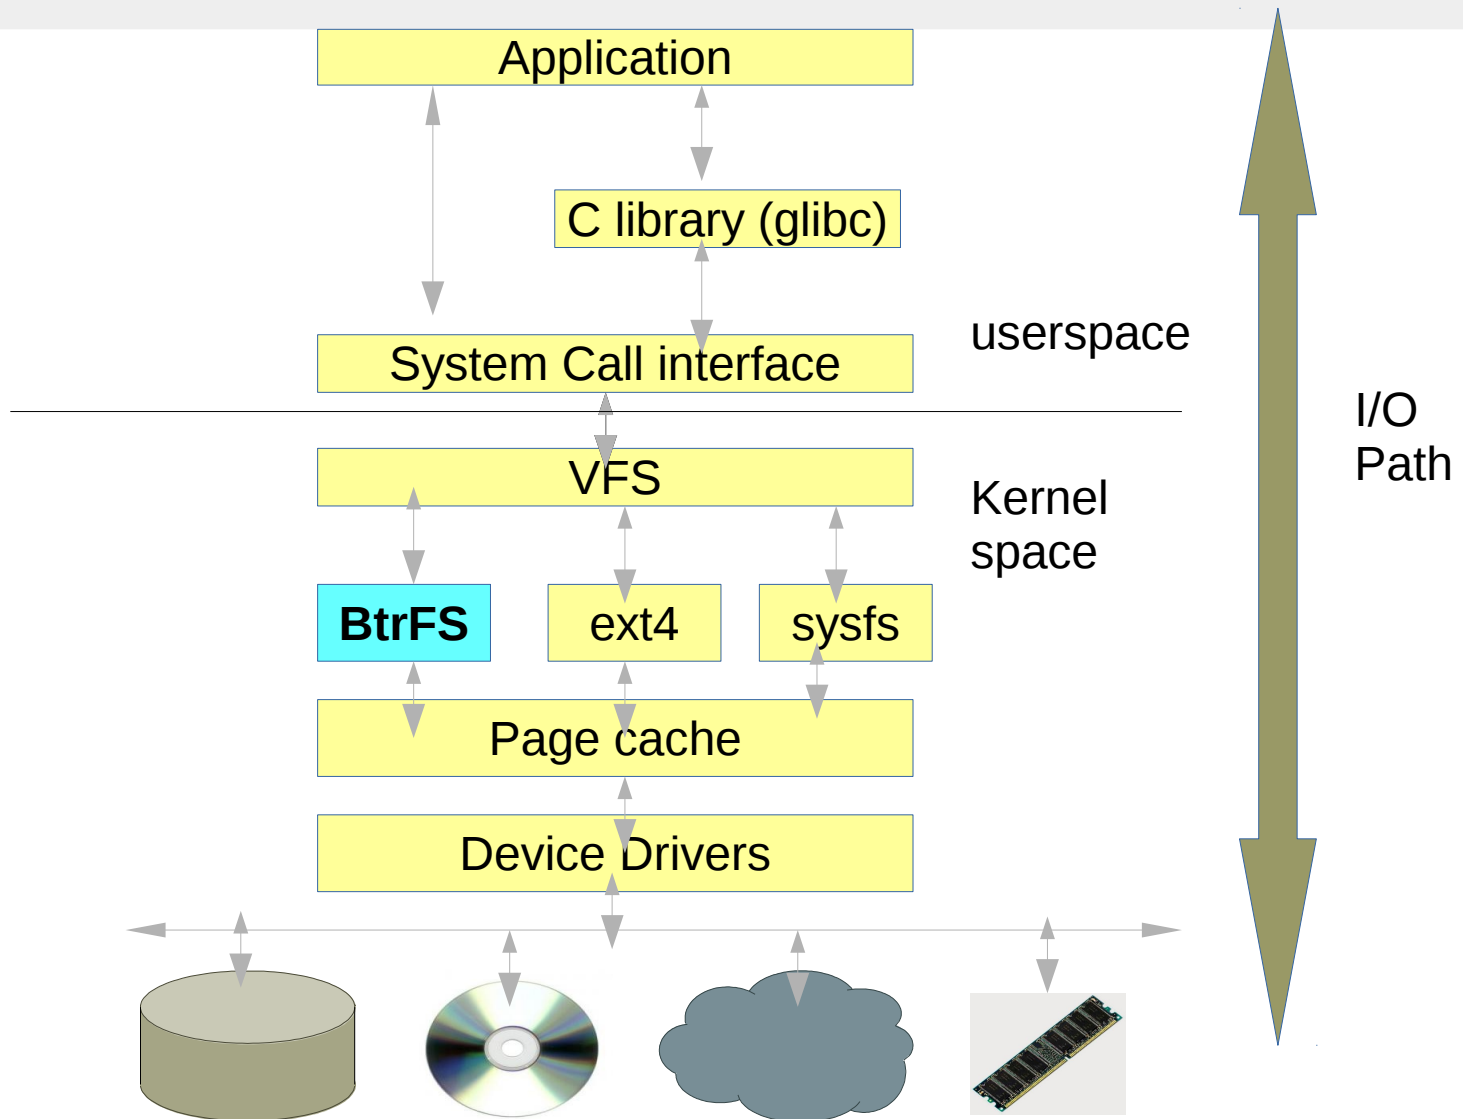
Hello everyone,

After the last FS summit, I started working on a new filesystem that maintains checksums of all file data and metadata. Many thanks to Zach Brown for his ideas, and to Dave Chinner for his help on benchmarking analysis.

outlines the Linux storage stack as of Kernel version 3.17



# Background – File System





# Background – File System

- Design Issue
  - Reliable Storage
    - Normal usage
    - Failure conditions
  - Fast Access
    - Different Scenarios
  - Efficient layout
    - Small files
    - Lots of files
- Operational Issues
  - Vulnerability windows
    - Log but only metadata
    - RAID write hole
  - Recovery
  - Defragmentation
  - Large directories
  - Resizing

Source: [dclug.tux.org/200908/BTRFS-DCLUG.pdf](http://dclug.tux.org/200908/BTRFS-DCLUG.pdf)

# Background – File System

- Design Issue
  - Reliable Storage
    - Normal usage
    - Failure conditions
  - Fast Access
    - Different Scenarios
  - Efficient layout
    - Small files
    - Lots of files
- Operational Issues
  - Vulnerability windows
    - Log but only metadata
    - RAID write hole
  - Recovery
  - Defragmentation
  - Large directories
  - Resizing

Source: [dclug.tux.org/200908/BTRFS-DCLUG.pdf](http://dclug.tux.org/200908/BTRFS-DCLUG.pdf)

# Background – File System

- Design Issue
  - Reliable Storage
    - Normal usage
    - Failure conditions
  - Fast Access
    - Different Scenarios
  - Efficient layout
    - Small files
    - Lots of files
- Operational Issues
  - Vulnerability windows
    - Log but only metadata
    - RAID write hole
  - Recovery
  - Defragmentation
  - Large directories
  - Resizing

Source: [dclug.tux.org/200908/BTRFS-DCLUG.pdf](http://dclug.tux.org/200908/BTRFS-DCLUG.pdf)

# Btrfs Features

- Key features
  - Extent based file system
  - Writable snapshots, read-only snapshots
  - Subvolumes (separate internal filesystem roots)
  - Checksums on data and metadata (crc32c)
  - Integrated multiple device support
    - File Striping, File Mirroring, File Striping+Mirroring, Striping with Single and Dual Parity implementations
  - Background scrub process
  - Offline filesystem check
  - Online filesystem defragmentation

# Main principles

- Inodes
- Extent based file system
- File system layout
- COW friendly B-trees
- Snapshot
- Software RAID
- Subvolume

# Inodes

- Everything in Linux is a file (Idea borrowed from Unix)
- Metadata (Data structure) associated with a file (Owner, permission, timestamp, actual file location(disk-blocks) and plethora of other information)  
**VFS layer representation:** struct inode
- Stored on a disk (Persistent storage)
- **BtrFS representation:** struct btrfs\_inode\_item

```
[amit@discworld papers]$ ls -li
total 7252
12853689 drwx-----. 3 amit amit    4096 Oct 26  2014 FAST
12722069 -rw-----. 1 amit amit  130652 Oct 26  2014 flash-enterprise.pdf
12722070 -rw-----. 1 amit amit 3421258 Oct 26  2014 jcse_3-3_50.pdf
12722071 -rw-----. 1 amit amit  117576 Oct 26  2014 novos-hotos2011.pdf
12722072 -rw-----. 1 amit amit 1580486 Oct 26  2014 NVM13-Wheeler_Linux_and_NVM.pdf
13369519 drwxr-xr-x. 2 amit amit    4096 Nov  1  2014 OS-papers
12722073 -rw-----. 1 amit amit   95059 Oct 26  2014 storage-topology.pdf
12722074 -rw-----. 1 amit amit   902215 Oct 26  2014 vol4no9main_part22.pdf
12722075 -rw-----. 1 amit amit 1155377 Oct 26  2014 xie.pdf
```

# Extent

- Extents are physically contiguous area on storage (Say disks).
- File consists of zero or more extents.
- I/O takes place in units of multiple blocks if storage is allocated in consecutive blocks. For sequential I/O, multiple block operations are considerably faster than block-at-a-time operations.
- Easier book-keeping


# COW

- Copy-On-Write
  - Technique
    - Every consumer is given pointer to the same resource.
    - Modification attempt leads to Trap
    - Create a local copy
    - Modify the local copy
    - Update the original resource
  - Benefit
    - Crash during the update procedure does not impact the original data.




# COW contd...

- Btrfs uses CoW style data/metadata update
  - Safer than overwrite style update by design
- Overwrite style: Update the data in place




- CoW style: Copy, update, and replace pointer



# COW friendly (B){ayer|oeing|balanced|broad}-tree

- Main idea
  - B+ -tree
  - Top down update procedure
  - Remove Leaf chaining
  - Lazy reference counting (Cloning purposes, Utilizing DAG)



(a) A basic b-tree (b) Inserting key 19, and creating a path of modified pages. [3]

# Data/Metadata checksum


- Btrfs has checksum for each data/metadata extent to detect and repair the broken data.
- Metadata is redundant and checksummed and data is checksummed too.
- When Btrfs reads a broken extent, it detects checksum inconsistency.
  - With mirroring: RAID1/RAID10
    - Read a correct copy
    - Repair a broken extent with a correct copy
  - Without mirroring
    - Dispose a broken extent and return EIO
- With “*btrfs scrub*”, Btrfs traverses all extents and fix incorrect ones
  - Online background job
- Demo

# Data/Metadata checksum



Original

# Data/Metadata checksum



After flipping **one** bit

# Data/Metadata checksum


- Demo
  - Corrupt disk block by bypassing file system to simulate disk faults.
  - Result
    - <snip from kernel logs>  
*BTRFS info (device sdb): csum failed ino 257 off 0 csum 2566472073 expected csum 3681334314*  
*BTRFS info (device sdb): csum failed ino 257 off 0 csum 2566472073 expected csum 3681334314*  
*BTRFS: read error corrected: ino 257 off 0 (dev /dev/sdc sector 449512)*
    - <snip from kernel logs>
      - BtrFS is able to identify such errors and corrects them too :)




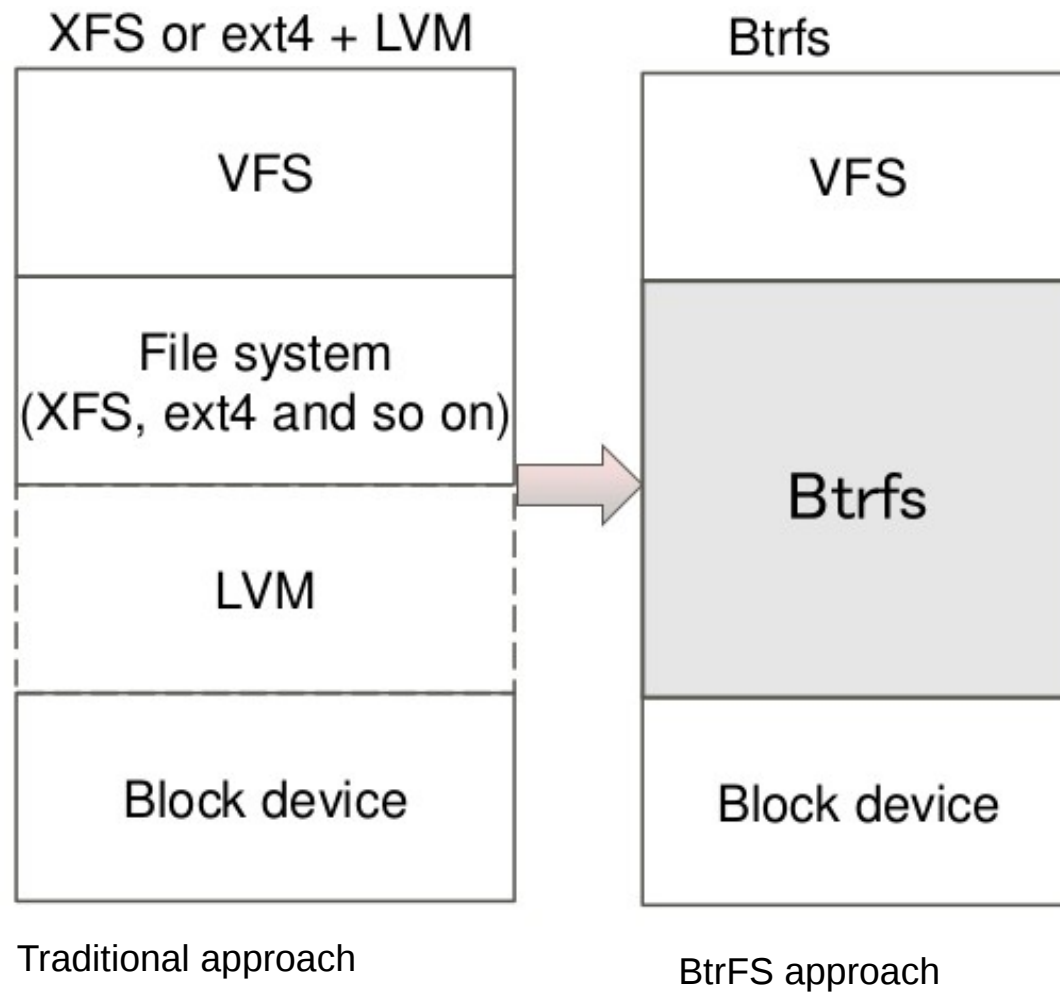
# Software RAID

- Motivation

- Increased capacity
- Greater reliability.
- Software modules that take raw disks, merge them into a virtually contiguous block address space, and export that abstraction to higher level kernel layers.
- Support mirroring, striping, and RAID5/6.



# Software RAID



Source: Linux foundation [8]



# Software RAID contd...

- Raid 0, 1, 5 (single parity) and 6 (redundant parity) are built in BtrFS
- Metadata is usually stored in duplicate form in Btrfs filesystems, even when a single drive is in use. (Also allows administrator to explicitly specify the raid level for metadata and data separately).
- Dynamically increase the storage
- Example

```
# mkfs.btrfs -d <mode> -m <mode> <dev1> <dev2> ...
```

```
# mount /dev/<diskX> /<mount-point>
```

– Where  $\text{diskX} \in \{\text{dev}\}$

```
# btrfs device add /dev/<diskY> /<mount-point>
```

# Scrub

- Btrfs CRCs allow to verify data stored on disk
- CRC errors can be corrected by reading a good copy of the block from another drive
- **ONLINE** background filesystem scrub (Not a full fsck)
- Scrubbing code scans the allocated data and metadata blocks
- Any CRC errors are fixed during the scan if a second copy exists

# Subvolumes

- Subvolumes allow the creation of multiple filesystems on a single device (or array of devices).
- Each subvolume can be treated as its own filesystem and mounted separately and exposed as needed (No need to mount the “root” device at all)
- Unlike subvolumes in LVM(an independent logical volume, composed of physical volumes), it has hierarchy and relations between subvolumes.
- Example
  - # btrfs subvolume create <user-defined-name>*

# Snapshot

- Works on subvolumes
- Snapshot is a cheap atomic copy of a subvolume, stored on the same file system as the original.
- Snapshots have useful backup function.
- Snapshot of snapshot of snapshot ...
- Read-only or a Read-write snapshot

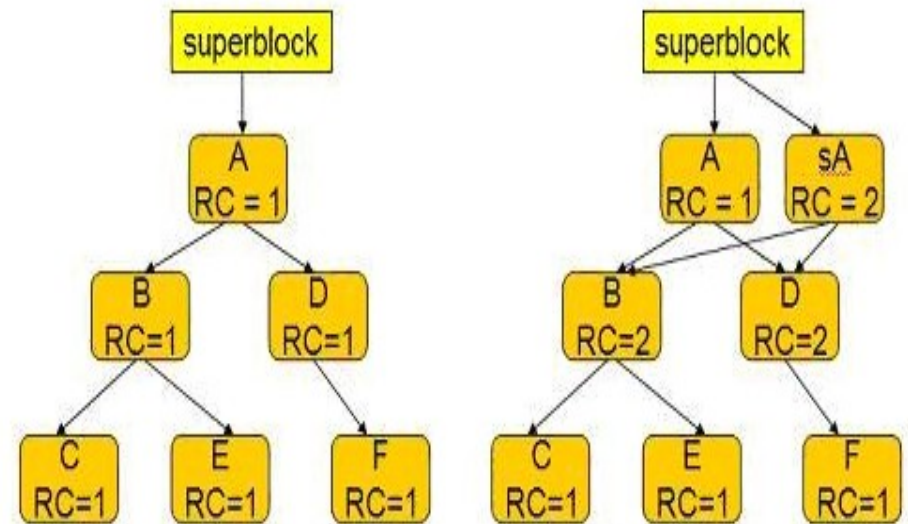


Fig: Snapshot of subvolume[4]

# Robustness

- **Test**

- Tolerance to Unexpected Power Failure while Writing to Files

- **Environment**

- Custom board (Freescale TWR-VF65GS10)
- Memory : 1GB DDR3
- Storage : 16GB Micro SD Card
- Software: Linux Kernel 3.15-rc7
- Application: Power Supply Control Unit Periodically Turns On and Off DC Power Supply every Minute, while a file writing application continuously Creates 4KB Files and Writes to it.
- Test candidates: Ext4, Btrfs

- **Result**

- Ext4 was corrupted and needed a file system check (fsck), while Btrfs didn't show any abnormalities.

Source: Fujitsu [5]

# Robustness

	Number of Power Failure	Results
Btrfs	1,000+	No Abnormal Situation Occurred
Ext4	1,000+	<b>Corrupted inode</b> had increased up to 32,000 and Finally Fell into Abnormal <b>Disk Full</b> State

Table: BtrFS vs Ext4 power failure test result

Source: Fujitsu [5]

# Performance

- **Test**

- Basic File I/O Throughput and Throughput under High Load

- **Environment**

- Intel Desktop Board D510MO
- Processor : 1.66 GHz Dual-Core Atom (4 Core with HT)
- Memory : 1GB DDR2-667 PC2-5300
- Storage : 32GB Intel X25-E e-SATA SSD
- Software: Linux Kernel 3.15.1 (x86\_64)
- Application: FIO (Single (for Basic) and Multiple (for High Load) FIO Running)
- Test candidates: Ext4, BtrFS

- **Result**

- I/O throughput under single FIO
  - Read : in Seq, Btrfs was Slightly Faster than Ext4
  - Write : Ext4 was almost Twice Faster than Btrfs
- I/O throughput under high load
  - Ext4: Every I/O Throughput Decreased Significantly, BtrFS: Decreased Less than Ext4

Source: Fujitsu [5]

# Performance

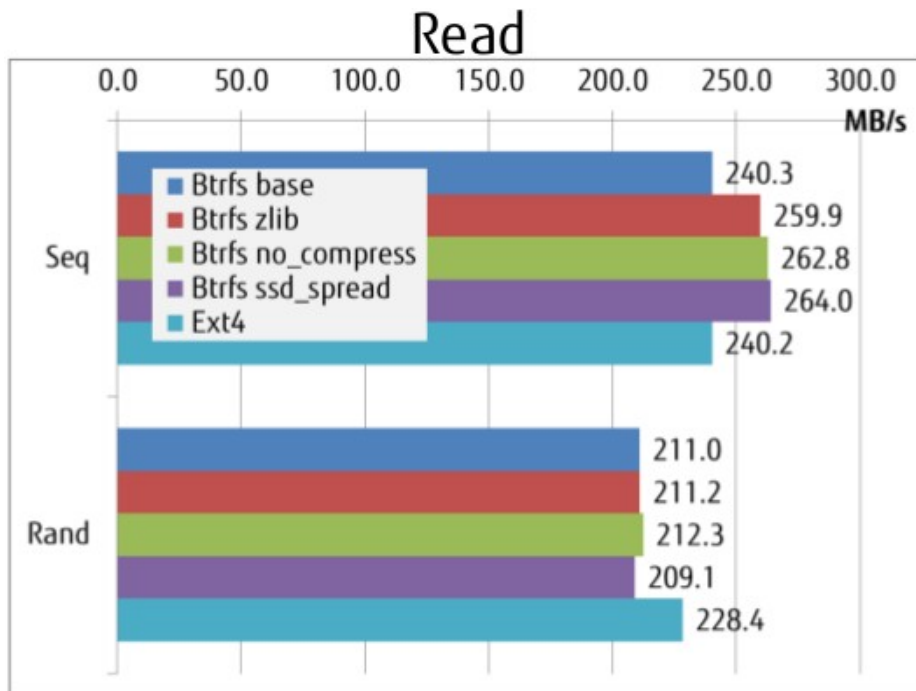


Fig: Read operation with single FIO

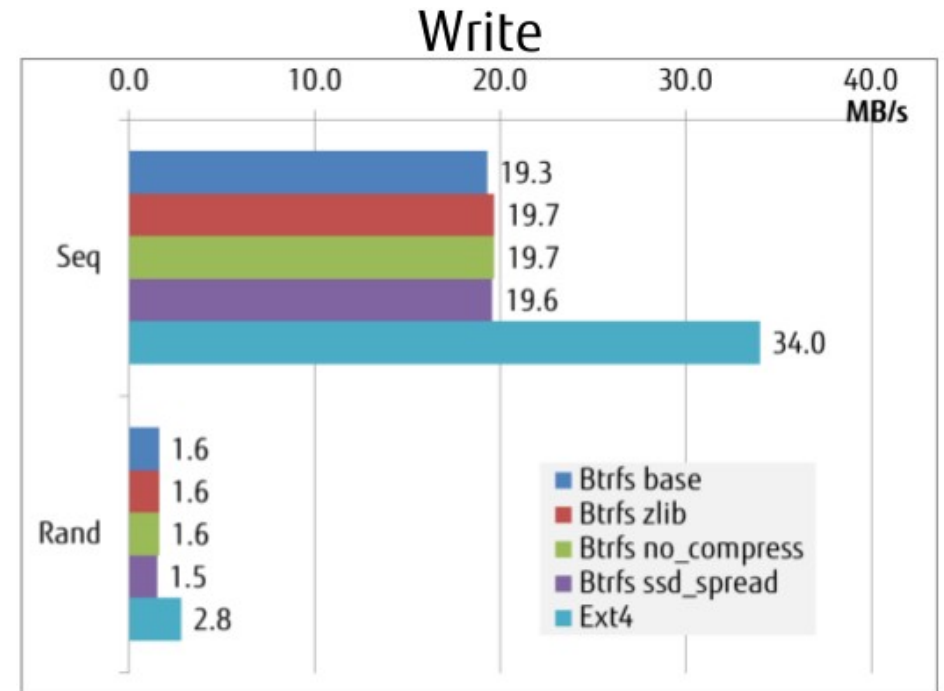



Fig: Write operation with single FIO

Source: Fujitsu [5]



# Who's using BtrFS



And many more...

# Gotchas

- Typically doesn't corrupt itself with latest kernel, but sometimes it does, So Always have backups.
- The Btrfs code base is under heavy development.
- Can get out of balance and require manual re-balancing.
- Auto de-fragmentation has problems with journal and virtual disk image files.
- Raid 5, 6 are still experimental.



# Conclusion

- Storage systems are not perfect, faults are inevitable.
- File systems play a crucial role in storage fault tolerance, data recovery, scalability and performance of the overall system.
- BtrFS[6] is a relatively new “copy-on-write” file system whose main design focus is fault tolerance.
- BtrFS adds an additional layer of storage fault tolerance capability to existing solutions and in some cases outperforms others or make them irrelevant.

# Reference

- [1] Ohad Rodeh. 2008. B-trees, shadowing, and clones. Trans. Storage 3, 4, Article 2 (February 2008), 27 pages. DOI=10.1145/1326542.1326544 <http://doi.acm.org/10.1145/1326542.1326544>
- [2] [http://www.phoronix.com/scan.php?page=news\\_item&px=mty0ndk](http://www.phoronix.com/scan.php?page=news_item&px=mty0ndk)
- [3] Ohad Rodeh, Josef Bacik, and Chris Mason. 2013. BTRFS: The Linux B-Tree Filesystem. Trans. Storage 9, 3, Article 9 (August 2013), 32 pages. DOI=10.1145/2501620.2501623 <http://doi.acm.org/10.1145/2501620.2501623>
- [4] <http://www.ibm.com/developerworks/cn/linux/l-cn-btrfs/>
- [5] [events.linuxfoundation.jp/sites/events/files/slides/linux\\_file\\_system\\_analysis\\_for\\_IVI\\_systems.pdf](http://events.linuxfoundation.jp/sites/events/files/slides/linux_file_system_analysis_for_IVI_systems.pdf)
- [6] <https://btrfs.wiki.kernel.org/>
- [7] Bianca Schroeder and Garth A. Gibson. 2007. Disk failures in the real world: what does an MTTF of 1,000,000 hours mean to you?. In Proceedings of the 5th USENIX conference on File and Storage Technologies (FAST '07). USENIX Association, Berkeley, CA, USA, Article 1
- [8] Satoru Takeuchi. Fujitsu. 2014. Btrfs Current Status and Future Prospects

# Be brave, try BtrFS TODAY

