

Linux Kernel IO subsystem

How it works and how can I see what is it doing?

Jan Kára

jack@suse.cz

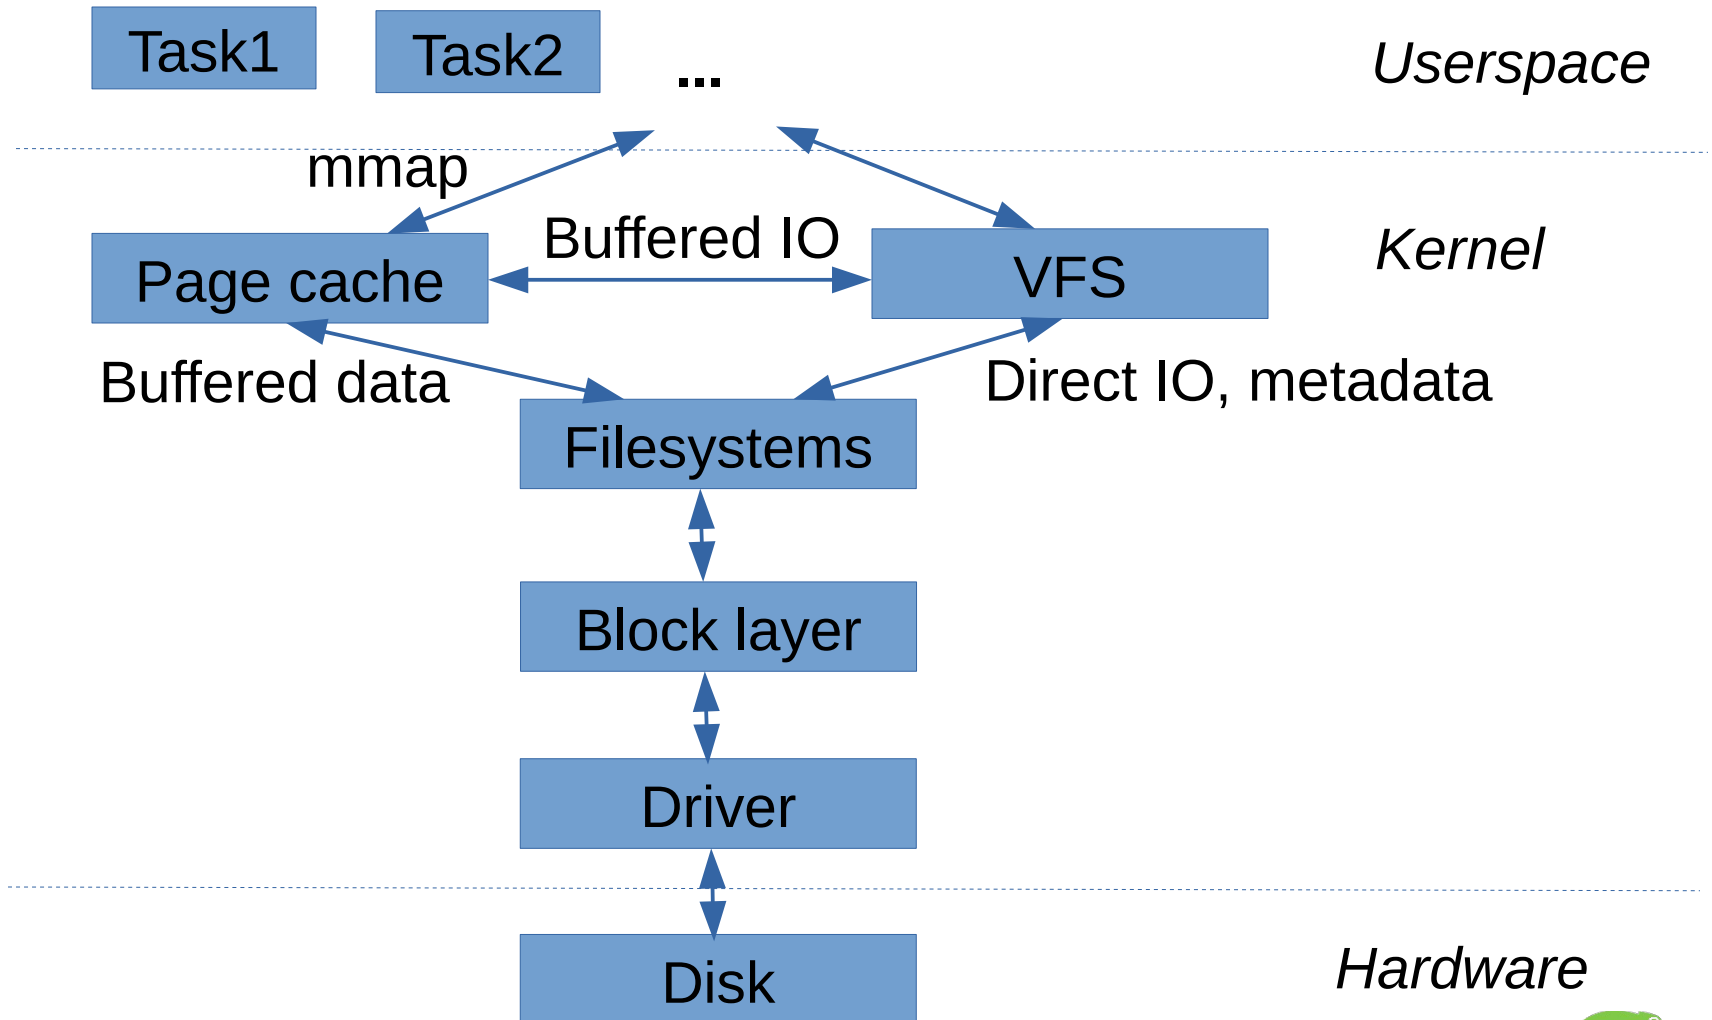


Outline

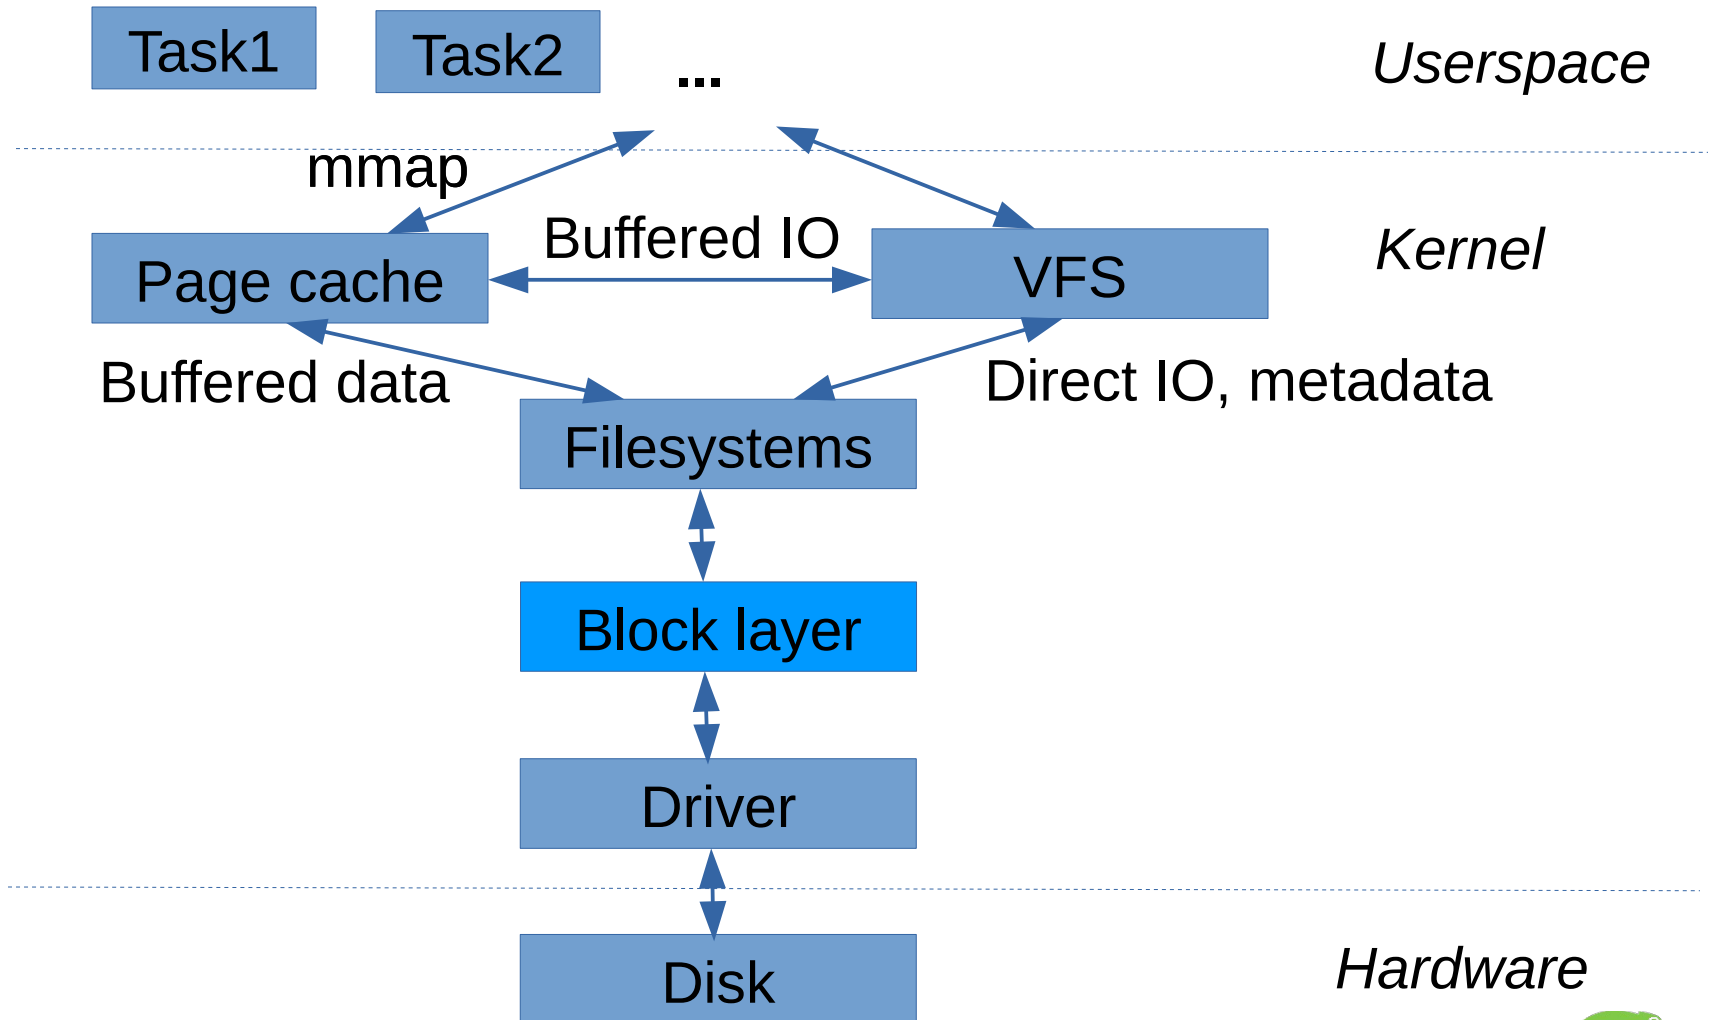
- Basic architecture of Linux IO subsystem
- Tools
- Examples

Linux Kernel IO Subsystem

Linux Kernel IO Architecture



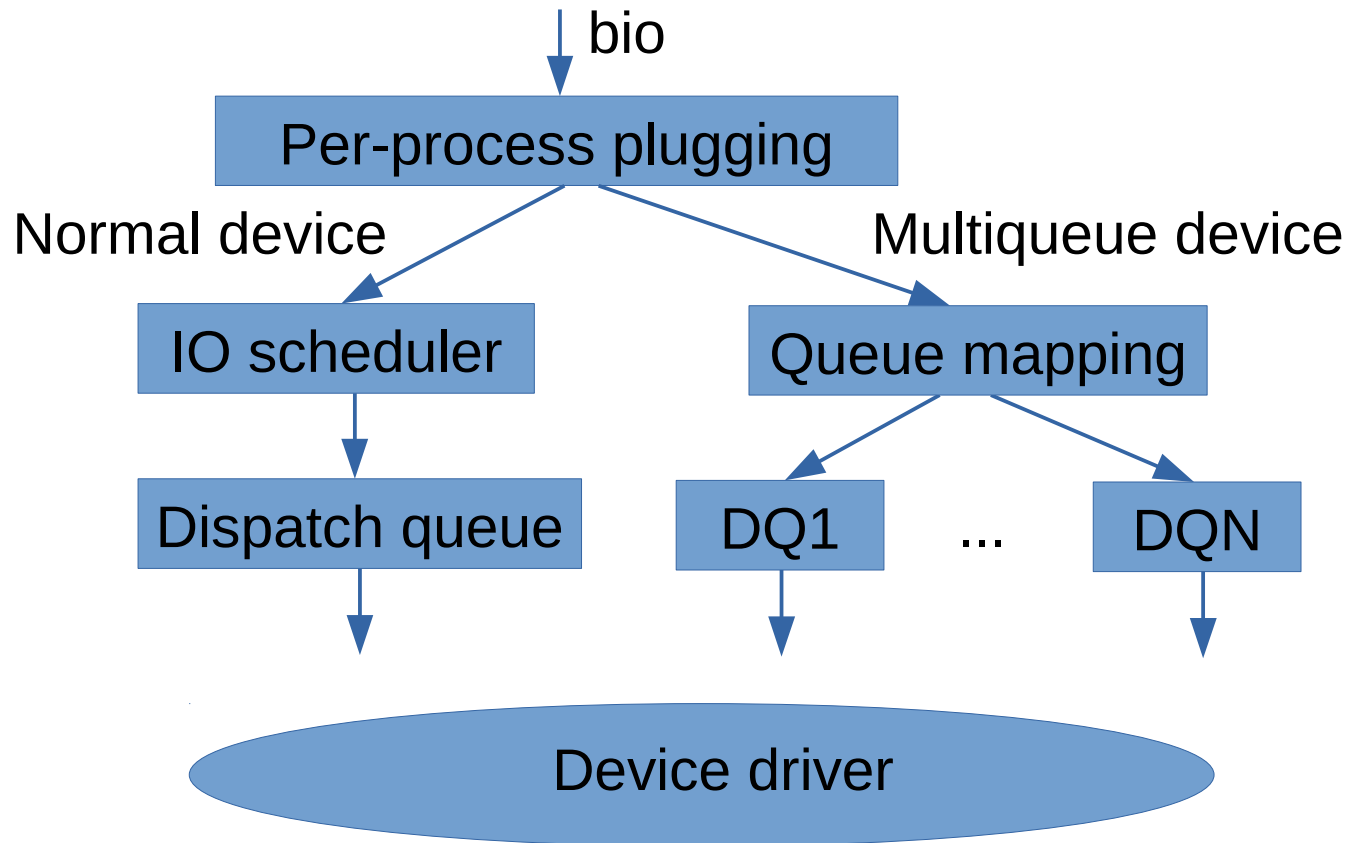
Linux Kernel IO Architecture



Block Layer Basics

- Works with IO requests
 - Starting sector, length, read / write / special
 - Can have hints (SYNC) and other flags (FUA, FLUSH)
- Life of a request
 - Created in block layer when IO submitted by a filesystem
 - Can be delayed, merged (IO scheduler, multiqueue handling)
 - Dispatched into a device driver
 - Completed when IO is finished

Submission Handling in Block Layer



IO Schedulers

- Decide when and in which order IO requests are submitted
 - NOOP – just pass requests into dispatch queue
 - Deadline
 - Prefers reads over writes
 - Sorts waiting requests to reduce seeking
 - Aims to dispatch each request at latest after its deadline has expired
 - CFQ
 - Prefers sync requests over async
 - Tries to achieve fairness among tasks
 - Support for IO priorities, cgroups, sync request idling, ...

Multiqueue Device Handling

- Used for fast devices
- Limited plugging
- No IO scheduling
- Support for multiple hardware IO queues
- Lightweight

IO Performance Analysis

iotat

usually packaged in sysstat package

- Shows basic statistics about IO
- Very lightweight
- Run: iostat -dxk 1

Dev:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	svctm
sda	182.00	0.00	353.00	0.00	2152.00	0.00	12.19	0.89	2.53	2.50
dm-0	0.00	0.00	536.00	0.00	2144.00	0.00	8.00	1.42	2.65	1.74

Dev:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	svctm
sda	235.00	0.00	300.00	0.00	2160.00	0.00	14.40	0.87	2.89	2.86
dm-0	0.00	0.00	540.00	0.00	2160.00	0.00	8.00	2.01	3.72	1.79

Example: Storage attached via Xen

- Sequential writes slow when writing through Xen blkfront driver to multipathed device
- Direct: 112 MB/s, through Xen: 46 MB/s
- iostat numbers:

Dev:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	svctm
Sdb	0.00	0.00	0.00	354.00	0.00	176128.00	995.07	31.97	91.79	2.84

IO through Xen

Dev:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	svctm
Sdd	0.00	0.00	0.00	1377.00	0.00	59988.00	87.13	30.98	22.38	0.73

After plugging fix (104 MB/s):

Dev:	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz	await	svctm
Sdb	0.00	0.00	0.00	571.00	0.00	145920.00	511.10	31.66	55.68	1.76



Takeaway 1

- Small requests hurt throughput
 - Overhead in kernel
 - Overhead in the device itself
- Holds for any storage device


iowatcher

<http://masoncoding.com/iowatcher/>

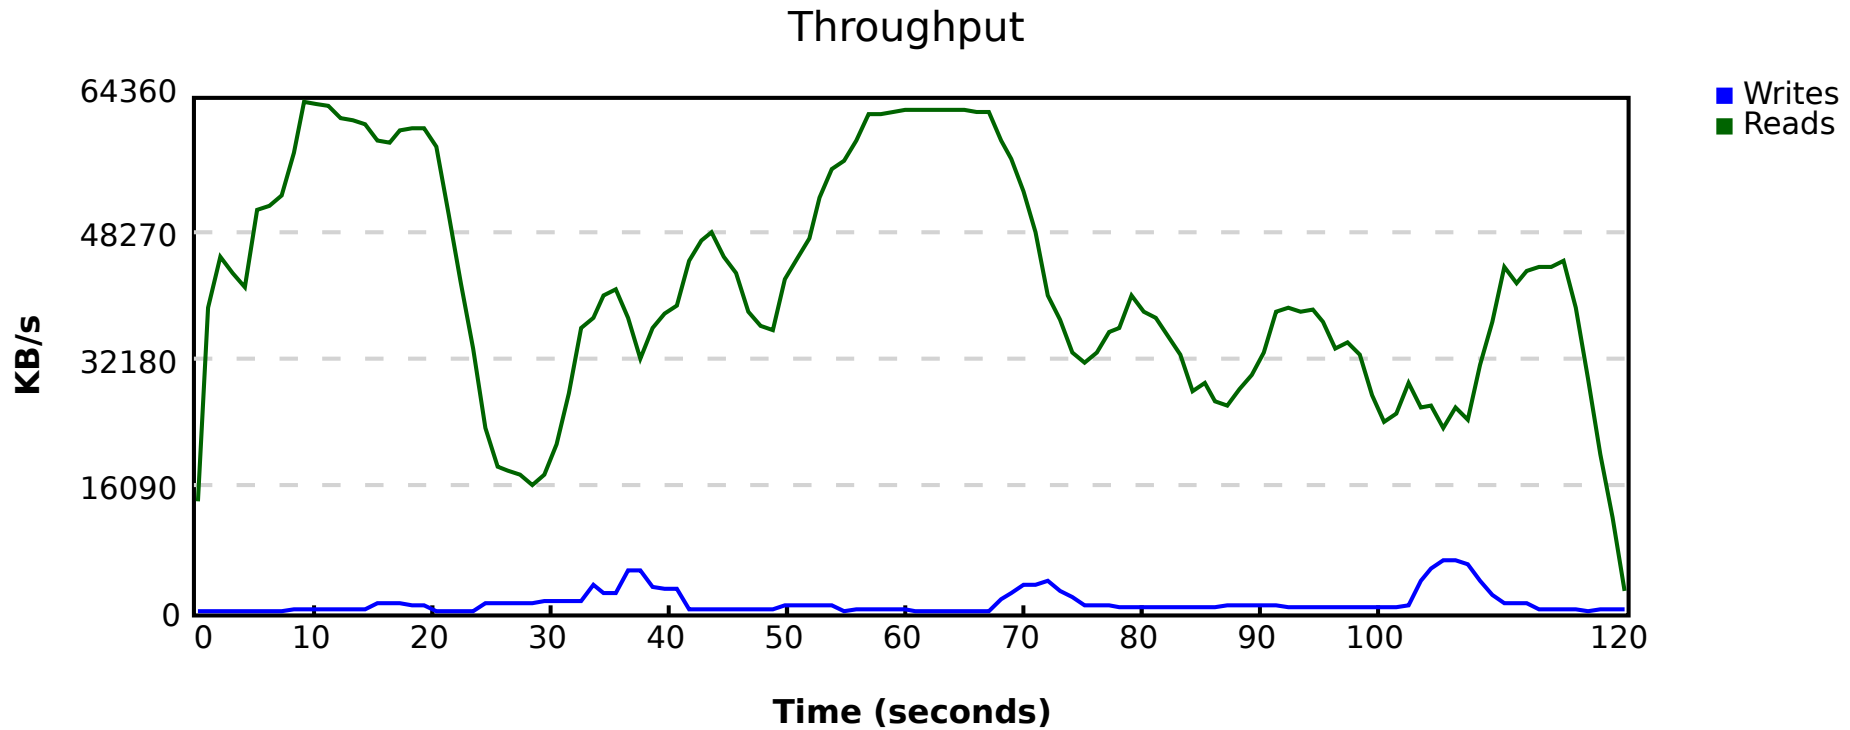
- Consumes traces from blktrace, mpstat output, fio bandwidth logs
- Plots various graphs (throughput, IO latency, IO location, ...)
- Movies

lowwatcher Graphs

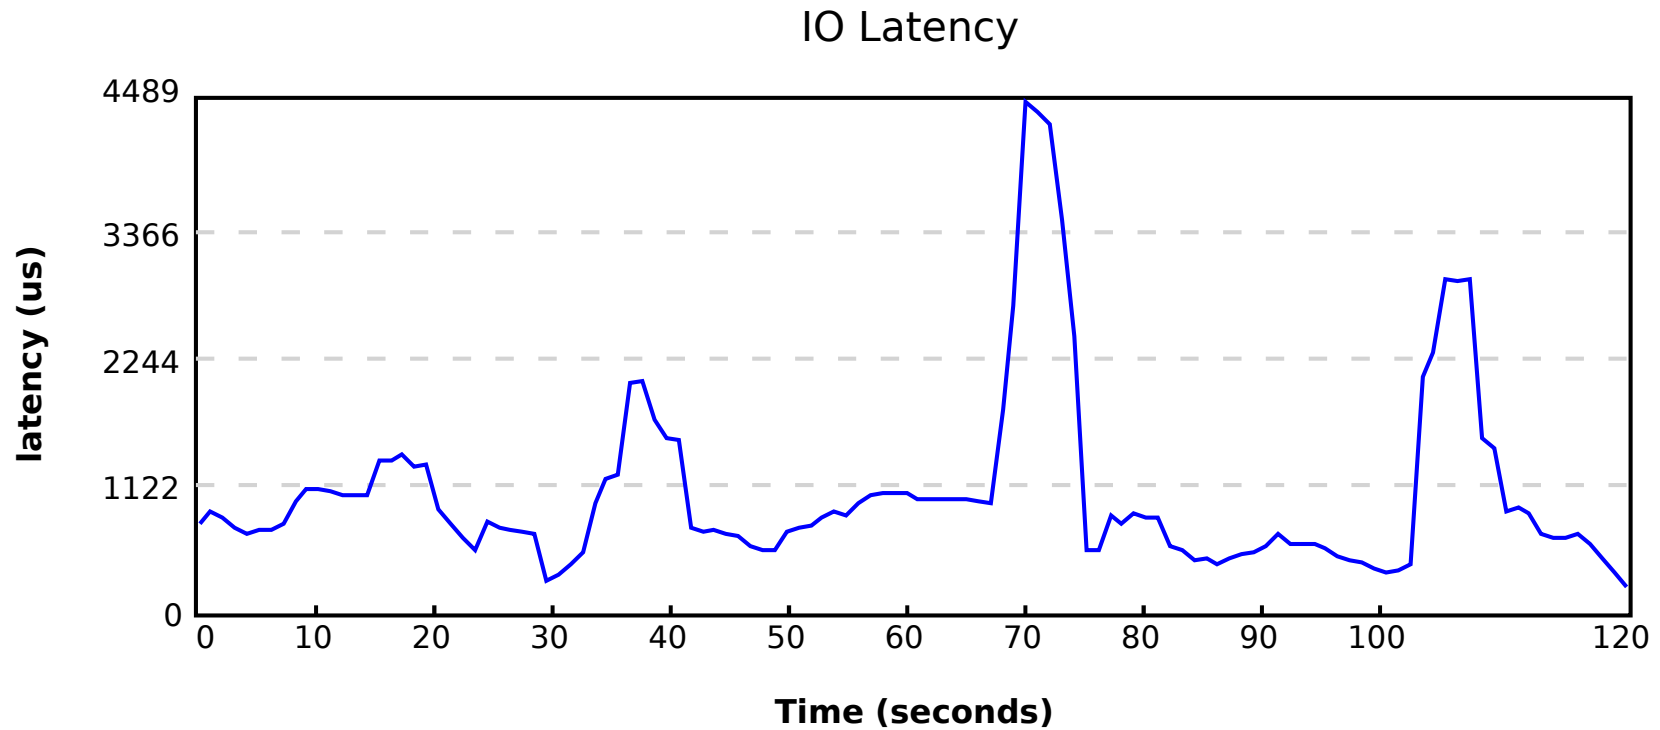
Device IO



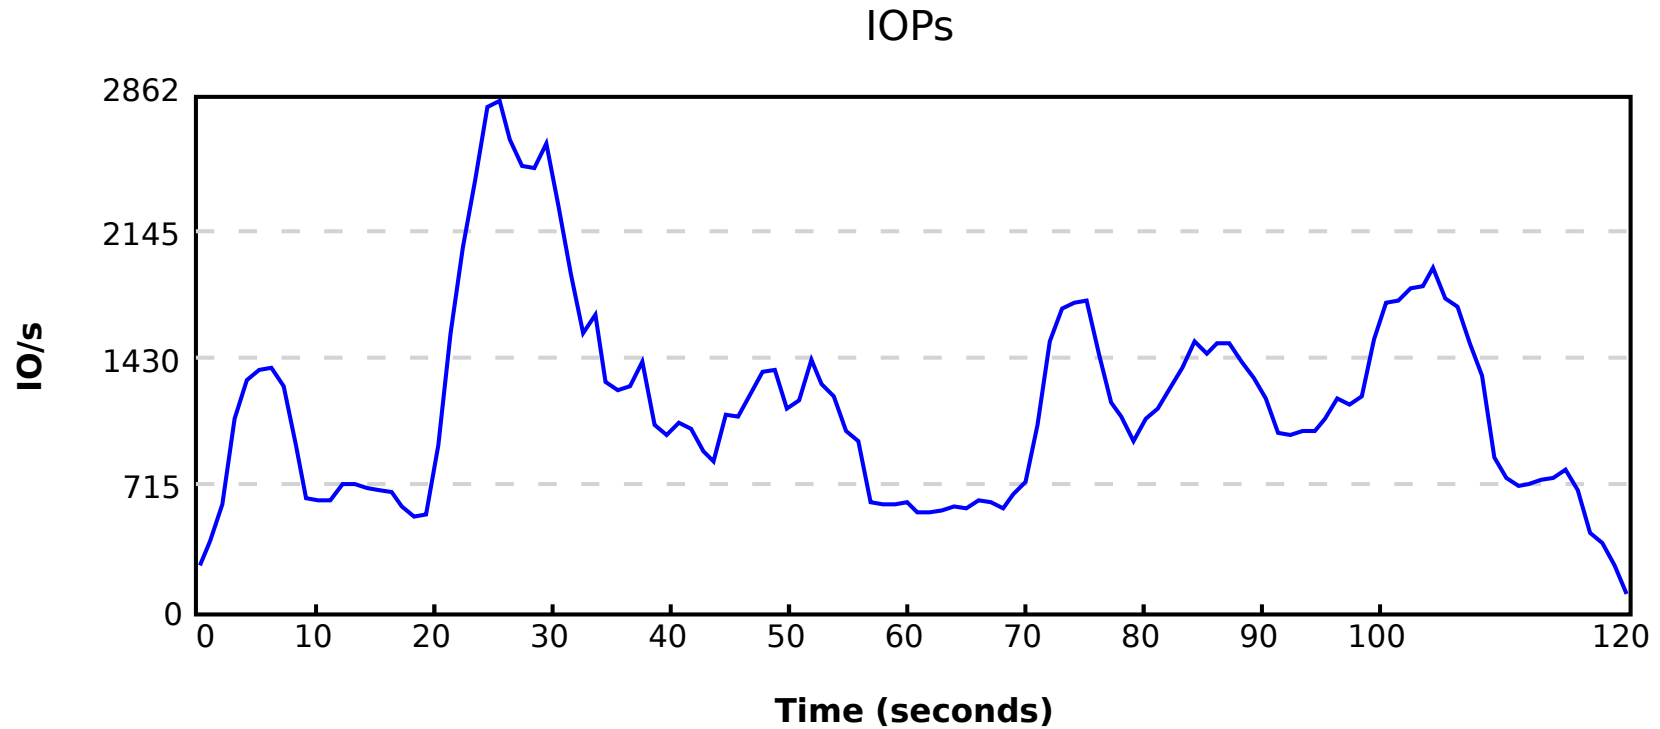
lowwatcher Graphs




lowwatcher Graphs



Iowatcher Graphs



lowwatcher Graphs



lowatcher Movie

- Movie time!

lowatcher Use Example

- Enabling ext4 dir_index feature slows down deleting of a directory tree.

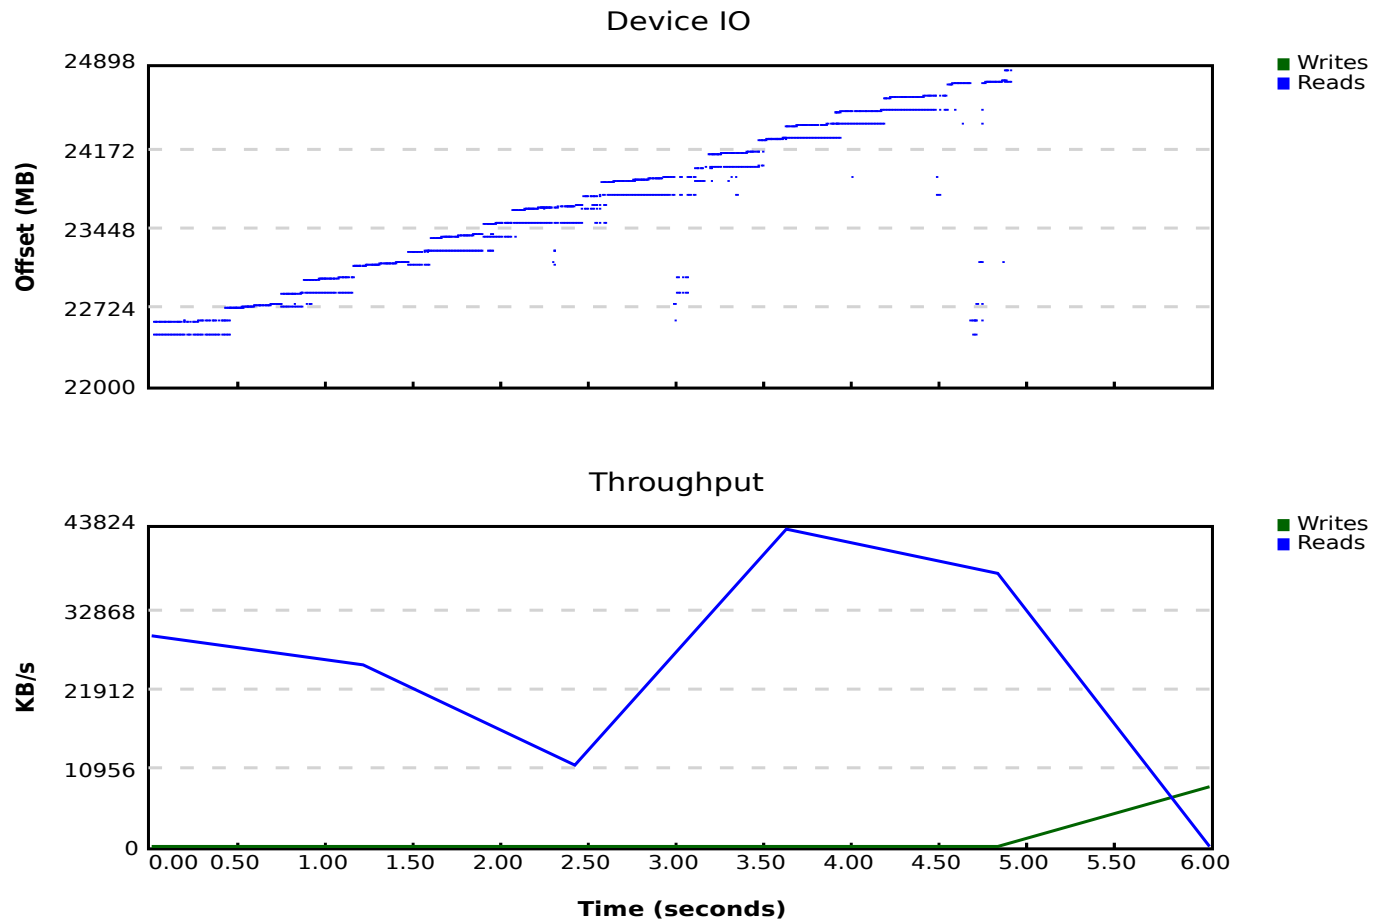
```
# Without dir_index
```

```
leet:~ # time rm -r /mnt/linux-2.6.32.orig/  
real 0m4.862s  
user 0m0.032s  
sys 0m2.388s
```

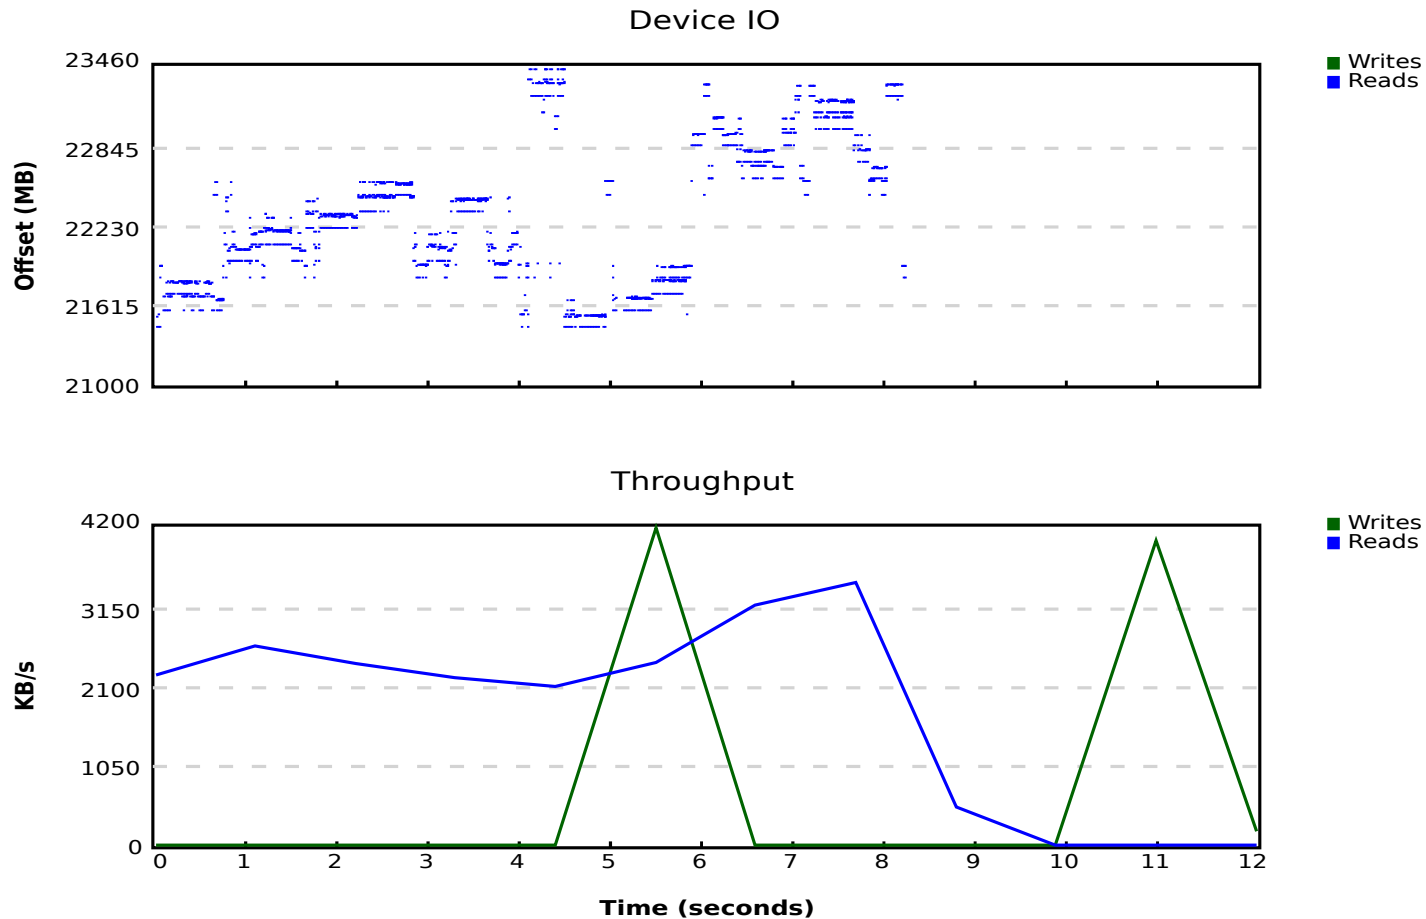
```
# With dir_index
```

```
leet:~ # time rm -r /mnt/linux-2.6.32.orig/  
real 0m8.100s  
user 0m0.040s  
sys 0m2.588s
```

Tree Deletion without Dir_index



Tree Deletion with Dir_index



Takeaway 2

- Seeking matters
 - Smaller chances of merging IO requests
 - Seek time for rotational storage

Blktrace

- Detailed information about IO requests processing
- Relatively large amount of data
 - May store over network
- Can handle multiple devices in parallel

- Gather as:

```
blktrace -d /dev/sda -d /dev/sdb -d /dev/dm-0
```

- Further processing using iowatcher, blkparse, btt

Blkparse

```
8,0 4 498 0.536245624 5072 A RM 46664392 + 8 <- (8,1) 46662344
8,0 4 499 0.536248072 5072 Q RM 46664392 + 8 [gcc]
8,0 4 0 0.536262021 0 m N cfq5072S / allocated
8,0 4 500 0.536262739 5072 G RM 46664392 + 8 [gcc]
8,0 4 501 0.536266614 5072 I RM 46664392 + 8 [gcc]
8,0 4 0 0.536268520 0 m N cfq5072S / insert_request
8,0 4 0 0.536270374 0 m N cfq5072S / add_to_rr
8,0 4 0 0.536276200 0 m N cfq workload slice:75
8,0 4 0 0.536278314 0 m N cfq5072S / set_active wl_prio:0
                        wl_type:2
8,0 4 0 0.536280939 0 m N cfq5072S / fifo=(null)
8,0 4 0 0.536282276 0 m N cfq5072S / dispatch_insert
8,0 4 0 0.536285224 0 m N cfq5072S / dispatched a request
8,0 4 0 0.536286509 0 m N cfq5072S / activate rq, drv=1
8,0 4 502 0.536286919 5072 D RM 46664392 + 8 [gcc]
8,0 4 503 0.556455119 0 C RM 46664392 + 8 [0]
8,0 4 0 0.556469202 0 m N cfq5072S / complete rqnoidle 0
8,0 4 0 0.556471881 0 m N cfq5072S / set_slice=25
8,0 4 0 0.556475942 0 m N cfq5072S / arm_idle: 2
                        group_idle: 0
8,0 4 0 0.556476510 0 m N cfq schedule dispatch
```

Blkparse (cont)

```

8,0 0 1 0.556659272 5073 A R 47002176 + 32 <- (8,1) 47000128
8,0 0 2 0.556661415 5073 Q R 47002176 + 32 [gcc]
8,0 0 0 0.556674617 0 m N cfq5073S / allocated
8,0 0 3 0.556675354 5073 G R 47002176 + 32 [gcc]
8,0 0 4 0.556677319 5073 P N [gcc]
8,0 0 5 0.556680380 5073 I R 47002176 + 32 [gcc]
8,0 0 0 0.556682649 0 m N cfq5073S / insert_request
8,0 0 0 0.556684273 0 m N cfq5073S / add_to_rr
8,0 0 6 0.556688402 5073 U N [gcc] 1
8,0 4 0 0.564839523 0 m N cfq idle timer fired
8,0 4 0 0.564842003 0 m N cfq5072S / slice expired t=0
8,0 4 0 0.564844568 0 m N / served: vt=9817282560
      min_vt=9817278464
8,0 4 0 0.564847483 0 m N cfq5072S / sl_used=2 disp=1
      charge=2 iops=0 sect=8
8,0 4 0 0.564849177 0 m N cfq5072S / del_from_rr
8,0 4 0 0.564850534 0 m N cfq schedule dispatch
8,0 4 0 0.564869775 0 m N cfq5073S / set_active wl_prio:0
      wl_type:2
8,0 4 0 0.564871692 0 m N cfq5073S / fifo=(null)
8,0 4 0 0.564872827 0 m N cfq5073S / dispatch_insert
8,0 4 0 0.564875317 0 m N cfq5073S / dispatched a request
8,0 4 0 0.564876421 0 m N cfq5073S / activate rq, drv=1
8,0 4 504 0.564876894 2743 D R 47002176 + 32 [kworker/4:2]
8,0 2 86 0.570193124 0 C R 47002176 + 32 [0]

```

Btt

- Uses binary event dump of blktrace events
- Produces various statistics
 - Q2C, Q2D, D2C latencies
 - Current device queue depth
 - Seeks per second
 - Per process activity
- Useful to check what to look for in blkparse output

Ftrace

- General kernel tracing framework
- Controlled via `/sys/kernel/debug/tracing`
- `Documentation/trace/ftrace.txt`
- Static trace points
 - Syscalls
 - Various events in filesystems, writeback, ...
- Dynamic trace points on almost every kernel function

Complex Problem Analysis

Problem

- When customer launches a large tarball creation, HA monitors of postgres DB occasionally time out ⇒ service failover
 - Service timeout 10 seconds
- Used HW raid for storage, deadline IO scheduler, ext3 filesystem
- 8 GB of memory free, disk is loaded with writes

Analysis 1/7

- iostat pretty normal:

Dev	rrqm/s	wrqm/s	r/s	w/s	rkB/s	wkB/s	avgrq-sz	avgqu-sz
Sda	0.00	13726.00	1.00	174.00	2.00	43876.00	501.46	142.12

- Blktrace output large (~900 MB parsed)
- Use btt to show latency outliers
 - Watch out for lost blktrace event

```
sort -k 2 -n -r q2clat.dat | head -30
```

```
127.842616 10.341348
```

```
127.842619 10.341346
```

```
127.842621 10.341344
```

```
...
```


Analysis 2/7

- Match back to blktrace events using timestamps

```
8,0 118.779433534 638 A WS 11705240 + 8 <- (8,2) 11395984
```

```
8,0 118.779433858 638 Q WS 11705240 + 8 [kjournald]
```

```
8,0 118.779435324 638 G WS 11705240 + 8 [kjournald]
```

```
8,0 118.779436253 638 I WS 11705240 + 8 [kjournald]
```

```
...
```

```
8,0 123.784506489 0 D WS 11705240 + 8 [swapper]
```

```
...
```

```
8,0 125.870800714 0 C WS 11705240 + 8 [0]
```

- Ok, so some IOs really take long
- Deadline IO scheduler delays outlying IO

Analysis 3/7

- Switched IO scheduler to NOOP
- Max latency reduced 3.8s and generally better
- Service time outs increased !?!
- Need more insight into why they time out
- Use ftrace to trace syscalls

```
echo 1 >/sys/kernel/debug/tracing/events/syscalls/enable  
cat /.../tracing/trace_pipe | gzip -c >syscall-trace.gz
```

Analysis 4/7

- Found large latencies in fsync

```
postgres-17461 [001] 3559.059091: sys_fsync(fd: 4)
```

```
postgres-17461 [008] 3570.848573: sys_fsync -> 0x0
```

```
syslog-ng-3008 [005] 3433.451593: sys_fsync(fd: 7)
```

```
syslog-ng-3008 [005] 3449.854534: sys_fsync -> 0x0
```

...

- Partly caused by heavy logging from sysrq-w
- Partly design limitation of data=ordered mode of ext3

Analysis 5/7

- Removed sysrq-w, switched ext3 to data=writeback
- Time outs less frequent but still occur
- Another syscall trace analysis

```
crm_master-20388 [000] 355206.448764: sys_read(...)
...
crm_node-20389 [006] 355207.654087: sys_mmap(...)
crm_node-20389 [006] 355207.654091: sys_mmap -> ...
crm_node-20389 [006] 355208.889691: sys_close(fd: 3)
crm_node-20389 [006] 355208.889693: sys_close -> 0x0
...
crm_master-20388 [001] 355220.880237: sys_read -> 0xa
```

Analysis 6/7

- Correlate inactivity periods with blktrace

```
8,0 283.784399307 1867 A R 9373112 + 8 <- (8,2) 9063856
8,0 283.784399608 1867 Q R 9373112 + 8 [crm_node]
8,0 283.784400643 1867 G R 9373112 + 8 [crm_node]
8,0 283.784401175 1867 P N [crm_node]
8,0 283.784401701 1867 I R 9373112 + 8 [crm_node]
8,0 283.784402232 1867 U N [crm_node] 1
...
8,0 284.987422579 0 D R 9373112 + 8 [swapper]
...
8,0 284.995404698 0 C R 9373112 + 8 [0]
```

Analysis 7/7

- Slow reads of shared libs sum up to time out
- Reads behind writes
- Switch back to deadline IO scheduler, tune deadlines and fifo_batch much lower
- Finally services run reliably

Conclusion

Conclusion

- Complex interactions between storage, block layer, filesystems
- Watch out for unexpected small requests, seeks
- Dependent reads vs async writes
 - Hard to guarantee latency under load

Thank You



Unpublished Work of SUSE LLC. All Rights Reserved.

This work is an unpublished work and contains confidential, proprietary and trade secret information of SUSE LLC. Access to this work is restricted to SUSE employees who have a need to know to perform tasks within the scope of their assignments. No part of this work may be practiced, performed, copied, distributed, revised, modified, translated, abridged, condensed, expanded, collected, or adapted without the prior written consent of SUSE. Any use or exploitation of this work without authorization could subject the perpetrator to criminal and civil liability.

General Disclaimer

This document is not to be construed as a promise by any participating company to develop, deliver, or market a product. It is not a commitment to deliver any material, code, or functionality, and should not be relied upon in making purchasing decisions. SUSE makes no representations or warranties with respect to the contents of this document, and specifically disclaims any express or implied warranties of merchantability or fitness for any particular purpose. The development, release, and timing of features or functionality described for SUSE products remains at the sole discretion of SUSE. Further, SUSE reserves the right to revise this document and to make changes to its content, at any time, without obligation to notify any person or entity of such revisions or changes. All SUSE marks referenced in this presentation are trademarks or registered trademarks of Novell, Inc. in the United States and other countries. All third-party trademarks are the property of their respective owners.

